

Hardness of Open-shop scheduling

Theorem: The problem $O3 \parallel C_{\max}$ (Minimize makespan of an open-shop schedule on 3 machines) is NP-hard.

Proof: Reduction from *Partition*. The input for *Partition* is a set S of n integers a_1, \dots, a_n whose total sum is $2B$. The problem is to find a subset of these integers whose total sum is exactly B . Given an instance S for *Partition* build the following instance for $O3 \parallel C_{\max}$:

There are $3n+1$ jobs, for the first n jobs (that is, for $j=1, \dots, n$) $p_{j,1}=a_j$, $p_{j,2}=p_{j,3}=0$. For the next n jobs (that is, for $j=n+1, \dots, 2n$) $p_{j,2}=a_j$, $p_{j,1}=p_{j,3}=0$, for the next n jobs (that is, for $j=2n+1, \dots, 3n$) $p_{j,3}=a_j$, $p_{j,1}=p_{j,2}=0$, and for the last job $p_{3n+1,1}=p_{3n+1,2}=p_{3n+1,3}=B$.

Claim: $C_{\max} = 3B$ if and only if the set S has a partition.

The proof has two parts – one for each direction of the ‘if and only if’

1. Assume there is a schedule with $C_{\max} = 3B$, it must be that the last job J_{3n+1} is processed on the machines one after the other with no idle. Assume w.l.o.g that M_2 is the second machine to process this job. Therefore, J_{3n+1} is processed on M_2 during the time interval $[B, 2B]$. As a result the processing of the n jobs $j=n+1, \dots, 2n$ with $p_{j,2}=a_j$, $p_{j,1}=p_{j,3}=0$ on M_2 splits between the intervals $[0, B]$ and $[2B, 3B]$, inducing a partition.

2. Assume the set S has a partition into two sets U and W each with total size B . A valid schedule with $C_{\max} = 3B$ is the following: On M_1 : J_{3n+1} followed by the first n jobs. On M_2 : the jobs of the second set of jobs ($j=n+1, \dots, 2n$) originated from U , then J_{3n+1} , then the jobs of the second set of jobs originated from W . On M_3 : the jobs $2n+1, \dots, 3n$ and then J_{3n+1} . Since a partition exists, the second machine is processed with no idle and the jobs from U and W exactly fill the B -segments before and after the processing of job J_{3n+1} .

Facility Location

Theorem: The network covering problem is NP-hard.

Proof: Reduction from Dominating Set.

A dominating set in an undirected graph is a collection S of vertices with the property that every vertex v in G is either in S , or there is an edge between a vertex in S and v .

Example: $o-x-o-o-x$

Given a graph G and an integer k , it is NP-hard to determine whether G has a DS of size k .

Given an instance of DS we build an instance of the covering problem: Same network, same k , all edges have length 1, all nodes have covering demand 1.

It is easy to see that a DS is exactly a cover.

Covering a tree

Theorem: The algorithm (slides 15-23) is optimal, that is, it uses the minimal possible number of centers.

Proof (draft): Let k be the number of centers determined by the algorithm. We show that there exists a set H of k nodes such that for every two nodes v_1, v_2 in H it holds that $d(v_1, v_2) > s_{v_1} + s_{v_2}$. (*)

Given that such a set exists, at least k centers are required, since otherwise, by the pigeonhole principle, there exists a center that covers more than one node in the set, which is impossible. The set H is defined as follows: For every center c set by the algorithm, add to H the node for which $d(v, c) = s_v$ and is the first node processed by the algorithm among the nodes covered by c .

Proof of the property (*): Consider a pair of vertices v_1, v_2 in H . W.l.o.g, the center covering v_1 was set first. Since G is a tree, there is only one path between v_1 and v_2 . By the way the algorithm proceeds, the center covering v_1 is located along this path (this requires more arguments from graph theory that we skip). Since it does not cover v_2 , it must be that that $d(v_1, v_2) > s_{v_1} + s_{v_2}$.